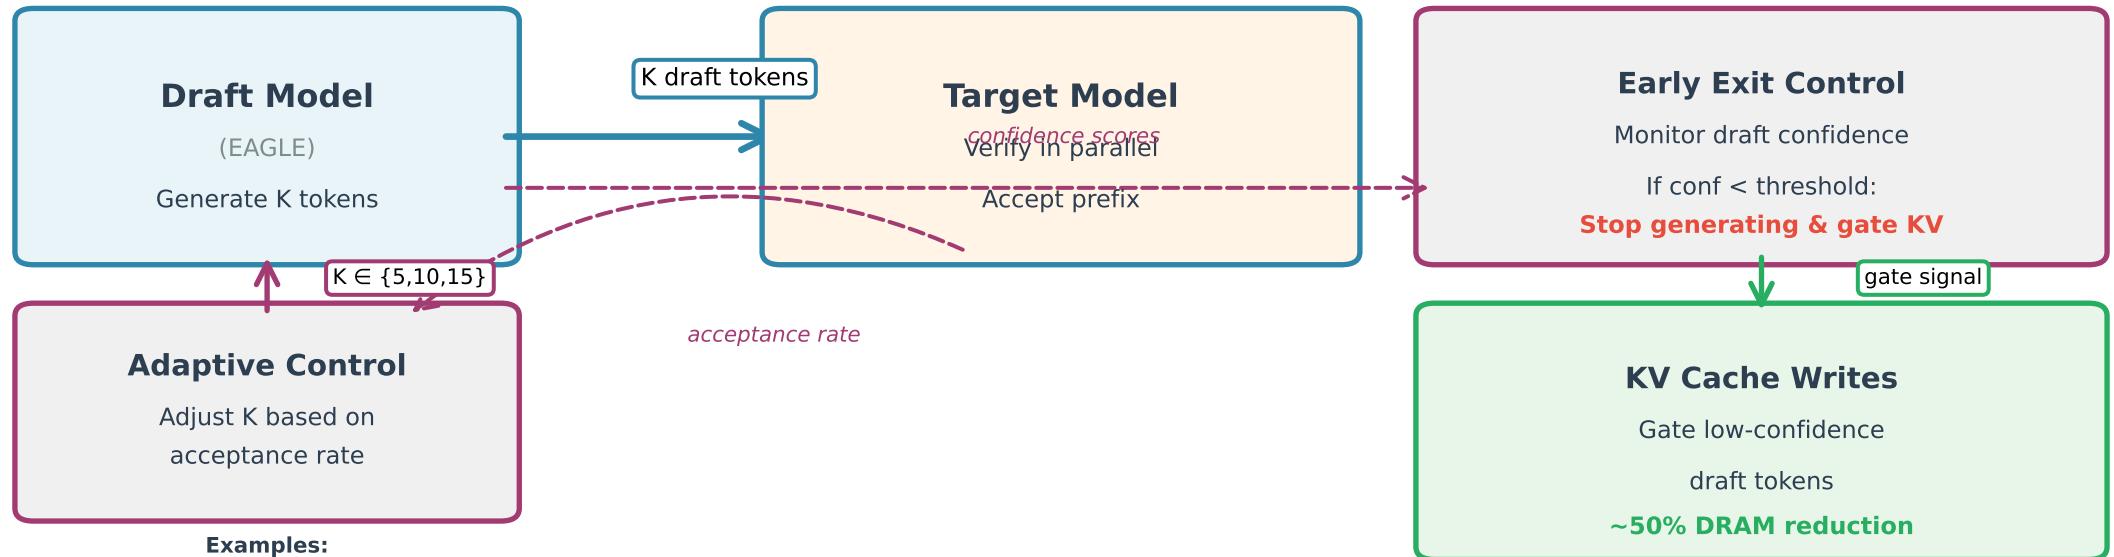


Elastic Speculation Overview

Adaptive Draft Length + Confidence-Based Early Exit (KV Gating)



□ Adaptive Draft Length

~30-50% lower latency vs. fixed-length

Dynamically adjusts speculation depth

□ Confidence-Based Early Exit

~50% less KV-cache DRAM traffic

Minimal latency impact (~1-3%)